

ARCHITECTURE RECOMMENDATION

robo_advisor

Conversational robo-advisor answering planning, tax, risk and product questions within the firm's methodology and policy, using the client's real portfolio through deterministic APIs and three knowledge bases; compliance-reviewed before answering.

Recommended architecture	Router / classifier dispatch (workflow)
Budget verdict	Exceeds the budget
Estimated p50 latency	~27 s of a 15 s budget
Estimated cost per request	\$0.414 at list prices, no caching
Tokens vs. one model call	2x (4 model calls)
Score	+1.1 (fit +4.2, latency -1.2, cost -1.9)
Currently implemented	Evaluator-optimizer loop (changes)

Recommended: **Router / classifier dispatch** (score +1.1; est. p50 latency ~27s vs budget 15s, ~2x tokens of one call). It exceeds the latency budget and no viable topology fits it: relax the budget, cut tool calls/reading per request, use a smaller model, or move work off the request path. The repository currently implements **Evaluator-optimizer loop** (score -5.7, ~82s); the recommendation changes the topology.

1. Why this architecture

- **+2.7** 4 distinct domains: route each request to a specialist prompt/agent instead of one prompt carrying every domain's instructions and tools. [OpenAI, A practical guide to building agents]
- **+1.5** Routing to specialists gives each a small, non-overlapping tool set. [OpenAI, A practical guide to building agents]

2. Repository scan

7 files scanned under /Users/saxena/Desktop/ClaudeCode_GIT/AgentTopologyBuilder/case_study/02_robo_advisor. Frameworks: anthropic-sdk, langgraph. Tools found: 10. Knowledge sources: elasticsearch. Side effects: external-post (irreversible). Current topology as implemented: **evaluator_optimizer** (Generator with an evaluator/critic and iteration cap.).

Current vs. recommended

	Current: Evaluator-optimizer loop	Recommended: Router / classifier dispatch
Score	-5.7	+1.1
Est. latency	~82 s	~27 s
Est. cost / request	\$0.673	\$0.414
Tokens vs. one call	2x	2x

Inferred profile fields

Field	Value	Conf.	Reason
tool_count	10	80%	10 tool definitions found (decorators, schemas, MCP tools, OpenAPI operations).
tool_calls_per_task	6	30%	Rough: 10 tools and 0 tool-handling call sites; verify against traces.
tool_dependency	sequential	40%	No concurrency primitives around tool calls.
tool_side_effects	irreversible	70%	Side-effecting integrations: external-post (irreversible)
knowledge_sources	1	60%	Retrieval integrations: elasticsearch
retrieval_depth	multi_hop	50%	Derived from sources (1) and whether a loop can issue follow-up queries.

Field	Value	Conf.	Reason
context_tokens_per_task	23000	30%	Rough function of sources and tools; replace with measured input tokens per request.
scope_breadth	1	40%	0 file(s) define system prompts/instructions; routing/handoffs raise this.
parallel_subtasks	1	40%	Single tool loop.
steps_predictable	False	60%	The model chooses steps at run time.
task_complexity	complex	40%	Tool loop with 10 tools.
latency_budget_s	8.0	40%	Request/response serving implies an interactive budget.
interaction	multi_turn	40%	Serving surface and memory usage.
verifiability	none	30%	No tests, evals or schemas found.
human_in_loop	True	60%	Approval/confirmation logic present.
error_recoverability	hard	50%	Irreversible integrations present.
accuracy_priority	4	40%	Irreversible actions raise the cost of errors.
tool_overlap	True	50%	Similarly named tools: search_methodology, search_policy, search_products

Evidence (first hits per category)

Category	Label	Location	Snippet
frameworks	anthropic-sdk	app/advisor.py:2	import anthropic
frameworks	anthropic-sdk	app/advisor.py:10	client = anthropic.Anthropic()
frameworks	langgraph	app/advisor.py:4	from langgraph.checkpoint.memory import MemorySaver
agent_patterns	tool runner	app/advisor.py:23	runner = client.beta.messages.tool_runner(
agent_patterns	tool runner	app/advisor.py:27	draft = runner.until_done()
agent_patterns	evaluator loop	app/advisor.py:1	"""Single-prompt advisor with a compliance reviewer pass."""
agent_patterns	evaluator loop	app/advisor.py:14	REVIEWER = """You are the compliance reviewer. Check the draft answer against policy and t
serving	http-api	app/advisor.py:9	app = FastAPI()
quality	human-approval	app/advisor.py:34	requires_confirmation = verdict.startswith("flag")
quality	memory	app/advisor.py:4	from langgraph.checkpoint.memory import MemorySaver
frameworks	anthropic-sdk	app/tools.py:3	from anthropic import beta_tool
tools	python decorator tool	app/tools.py:9	get_portfolio
tools	python decorator tool	app/tools.py:15	risk_score
tools	python decorator tool	app/tools.py:21	tax_lots
tools	python decorator tool	app/tools.py:27	contribution_limits
tools	python decorator tool	app/tools.py:33	suitability_check
retrieval	elasticsearch	app/tools.py:4	from opensearchpy import OpenSearch
retrieval	elasticsearch	app/tools.py:6	kb = OpenSearch(hosts=["https://kb.internal:9200"])
side_effects	external-post (irreversible)	app/tools.py:36	return requests.post("https://rules.internal/suitability", json={"client_id": client_id, "
side_effects	external-post (irreversible)	app/tools.py:66	return requests.post("https://proposals.internal/draft", json={"client_id": client_id, "su

3. Options compared

Score = structural fit + latency adjustment + cost adjustment. Latency is the estimated p50 on the critical path; cost sums every call at list prices. Non-viable options cannot do the job and are listed last.

#	Architecture	Score	Fit	Latency	Budget	Cost/req	Tokens	Calls
1	Router / classifier dispatch	+1.1	+4.2	~27 s	exceeds	\$0.414	2x	4
2	Parallel sectioning (code-defined fan-out)	-0.1	+0.7	~20 s	exceeds	\$0.045	1x	5
3	Orchestrator + worker subagents	-1.3	+1.5	~46 s	far over	\$0.104	2x	10
4	Single agent with tools	-1.6	+1.5	~26 s	exceeds	\$0.411	2x	3
5	Specialist handoffs (decentralised)	-3.9	+0.0	~38 s	far over	\$0.431	2x	5
6	Parallel voting / ensemble	-4.6	+0.0	~28 s	exceeds	\$2.067	8x	15
7	Prompt chain (sequential workflow)	-4.6	-2.0	~40 s	far over	\$0.097	2x	6
8	Evaluator-optimizer loop	-5.7	-1.0	~82 s	far over	\$0.673	2x	10
9	Hierarchical teams (orchestrator -> leads -> workers)	-11.2	-7.0	~74 s	far over	\$0.404	3x	26
10	Single model call not viable: needs tool calls; a single call cannot act	-5.0	-4.0	~9 s	fits	\$0.250	1x	1


Latency/accuracy frontier (no option is both faster and a better structural fit): Parallel sectioning (code-defined fan-out) (~20 s, fit +0.7), Single agent with tools (~26 s, fit +1.5), Router / classifier dispatch (~27 s, fit +4.2). Fastest viable within budget: **none**. Strongest structural fit: **Router / classifier dispatch**.

Runners-up

- **Parallel sectioning (code-defined fan-out)** (score -0.1, ~20 s, 1x tokens): Some of the fan-out can still be code-defined if a stable first split exists.
- **Orchestrator + worker subagents** (score -1.3, ~46 s, 2x tokens): Several domains inside one request: a lead delegates to domain specialists with narrow prompts and tool sets; ~2 independent sub-tasks discovered at run time: orchestrator spawns parallel subagents, each in its own context.
- **Single agent with tools** (score -1.6, ~26 s, 2x tokens): Unpredictable path: the model must choose steps at run time, which is the definition of an agent; Multi-hop retrieval: the model must decide follow-up queries from earlier results (agentic retrieval).

4. Decomposition plan: Router / classifier dispatch

A cheap classification step assigns each request to one of several specialised prompts, agents or models.



Component	Role	Model	Effort	Tools	Parallel
Router	Classifies the request into one category (small model, structured output).	claude-haiku-4-5	n/a	-	-
financial planning specialist	Handles 'financial planning' requests with only that domain's instructions and tools.	acme-onprem-70b	n/a	contribution_limits, create_proposal, get_portfolio, market_data, risk_score, search_methodology, search_policy, search_products, suitability_check, tax_lots	-
tax specialist	Handles 'tax' requests with only that domain's instructions and tools.	acme-onprem-70b	n/a	contribution_limits, create_proposal, get_portfolio, market_data, risk_score, search_methodology, search_policy, search_products, suitability_check, tax_lots	-
risk specialist	Handles 'risk' requests with only that domain's instructions and tools.	acme-onprem-70b	n/a	contribution_limits, create_proposal, get_portfolio, market_data, risk_score, search_methodology, search_policy, search_products, suitability_check, tax_lots	-
products specialist	Handles 'products' requests with only that domain's instructions and tools.	acme-onprem-70b	n/a	contribution_limits, create_proposal, get_portfolio, market_data, risk_score, search_methodology, search_policy, search_products, suitability_check, tax_lots	-

5. Model selection per role

Catalog case-study-ecosystem; data class **confidential**; verified entries only. Providers used: anthropic, self_hosted. Estimated per request on the chosen models: ~16 s, \$0.023.

Role	Needs	Chosen	Effort	Fallback	Est./call	Alternatives
router router	tier ≥2; 2s; 4,500 ctx; structured	claude-haiku-4-5	n/a	claude-sonnet-5	~0.8s \$0.0046	claude-sonnet-5 (+1.2), claude-opus-5 (+0.4)
handler_1 handler	tier ≥5; 13s; 49,500 ctx; tools	acme-onprem-70b	n/a	claude-sonnet-5	~10.0s \$0.0376	claude-sonnet-5 (+0.6), claude-opus-5 (-1.1)

Role	Needs	Chosen	Effort	Fallback	Est./call	Alternatives
handler_2 handler	tier ≥5; 13s; 49,500 ctx; tools	acme-onprem-70b	n/a	claude-sonnet-5	~10.0s \$0.0376	claude-sonnet -5 (+0.6), claude-opus-5 (-1.1)
handler_3 handler	tier ≥5; 13s; 49,500 ctx; tools	acme-onprem-70b	n/a	claude-sonnet-5	~10.0s \$0.0376	claude-sonnet -5 (+0.6), claude-opus-5 (-1.1)
handler_4 handler	tier ≥5; 13s; 49,500 ctx; tools	acme-onprem-70b	n/a	claude-sonnet-5	~10.0s \$0.0376	claude-sonnet -5 (+0.6), claude-opus-5 (-1.1)

- **router:** Tier 3 vs required 2; est. 0.8s of a 1.5s share; \$0.0046 per call. measured classification accuracy 93%
- **handler_1:** Tier 4 vs required 5; est. 10.0s of a 12.8s share; \$0.0376 per call. measured conversation accuracy 90% Fallback claude-sonnet-5 is on a different provider for outage/refusal resilience.
- **handler_2:** Tier 4 vs required 5; est. 10.0s of a 12.8s share; \$0.0376 per call. measured conversation accuracy 90% Fallback claude-sonnet-5 is on a different provider for outage/refusal resilience.
- **handler_3:** Tier 4 vs required 5; est. 10.0s of a 12.8s share; \$0.0376 per call. measured conversation accuracy 90% Fallback claude-sonnet-5 is on a different provider for outage/refusal resilience.
- **handler_4:** Tier 4 vs required 5; est. 10.0s of a 12.8s share; \$0.0376 per call. measured conversation accuracy 90% Fallback claude-sonnet-5 is on a different provider for outage/refusal resilience.

6. Cross-cutting recommendations

- **Build a ~20-query eval before adding any structure.** Every source agrees: measure a single agent/call first, then add topology only where the eval shows it failing. Above a ~45% single-agent baseline, adding agents tended to hurt in the scaling study. Keep the eval to compare latency and accuracy across topologies. [Google/MIT, Towards a Science of Scaling Agent Systems]
- **Prompt caching on the stable prefix (system prompt, tool schemas, reference docs).** Multi-turn and multi-call topologies re-send the same prefix on every call; caching cuts cost up to ~90% and time-to-first-token. [Anthropic, Effective context engineering for AI agents]
- **Structured outputs (JSON schema) for routers, extractors and votes.** Removes parsing failures and makes code-side gates and majority votes trivial. [Anthropic, Building effective agents]
- **Parallel tool use inside each agent turn.** Independent tool calls issued in one assistant turn cut rounds by 2-4x; return all results in one message. [Anthropic, Effective context engineering for AI agents]
- **Context editing (clear stale tool results/thinking).** Keeps long loops lean without summarisation losses. [Anthropic, Effective context engineering for AI agents]
- **Memory across sessions (file/DB-backed).** State that must outlive one conversation should be written to memory, not carried in context. [Anthropic, Effective context engineering for AI agents]
- **Dedicated, gated tools for irreversible actions with human approval.** Promote side-effecting actions to typed tools so the harness can intercept, confirm, audit and roll back; never let them run through a generic shell tool. [Anthropic, Building effective agents]
- **Model tiering: large model for planning/synthesis, smaller models for reading/routing.** Reading-heavy workers and routers need many input tokens and little judgment; run them on Sonnet/Haiku and keep Opus for the lead. [Anthropic, Managed Agents multiagent guidance]
- **Agentic retrieval capped at ~4 iterations: start broad, then narrow.** Multi-hop retrieval buys accuracy (27% -> 43% EM on HotpotQA) at 3-20x latency; cap the loop and prefer short broad queries first, narrowing progressively. [Agentic vs classic RAG guidance (see docs/industry_guidance.html)]

7. Estimate assumptions

- Small-model router (~20 output tokens) then one specialist handler.
- Serving assumptions (reference model per capability level from the catalog): claude-opus-5 ~1.5s TTFT, ~55 tok/s; acme-onprem-70b ~0.5s TTFT, ~90 tok/s; claude-haiku-4-5 ~0.6s TTFT, ~140 tok/s; tool call ~0.4s.
- Reading load per task 45,000 tokens (context pressure: medium); output type short_answer.
- Latency is the critical path (parallel branches count once); cost sums every call at list prices without caching. Treat both as order-of-magnitude.

Appendix A. Workload profile used

Field	Value	Field	Value
name	robo_advisor	task_complexity	complex
steps_predictable	False	scope_breadth	4
parallel_subtasks	2	output_type	short_answer
tool_count	10	tool_calls_per_task	5
tool_dependency	independent	tool_side_effects	irreversible
tool_overlap	True	tool_latency_s	0.4
knowledge_sources	1	retrieval_depth	multi_hop
context_tokens_per_task	45000	accuracy_priority	5
single_agent_baseline	0.61	verifiability	partial
error_recoverability	moderate	latency_budget_s	15.0
interaction	multi_turn	requests_per_day	8000
cost_sensitivity	3	human_in_loop	True
data_sensitivity	confidential	tools	['contribution_limits', 'create_proposal', 'get_portfolio', 'market_data', 'risk_score', 'search_methodology', 'search_policy', 'search_products', 'suitability_check', 'tax_lots']
sources	['methodology KB', 'policy KB', 'product KB']	domains	['financial planning', 'tax', 'risk', 'products']

Appendix B. Sources

- **ANTHROPIC_BEA**: Anthropic, Building effective agents
- **ANTHROPIC_CTX**: Anthropic, Effective context engineering for AI agents
- **ANTHROPIC_MA**: Anthropic, Managed Agents multiagent guidance
- **ANTHROPIC_MAR**: Anthropic, How we built our multi-agent research system
- **OPENAI_PG**: OpenAI, A practical guide to building agents
- **RAG**: Agentic vs classic RAG guidance (see docs/industry_guidance.html)
- **SCALING**: Google/MIT, Towards a Science of Scaling Agent Systems

Full evidence and the consolidated decision framework: docs/industry_guidance.html and docs/METHODOLOGY.html in the VG Select: 3A repository.